

```
# Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
# Load the dataset
data = pd.read_csv("customer_churn_dataset-testing-master[1].csv")

# Display first few rows
print("First 5 rows of dataset:")
print(data.head())

# Check for missing values and data types
print("\nDataset Info:")
print(data.info())

# Summary statistics
print("\nSummary Statistics:")
print(data.describe())
```

```
First 5 rows of dataset:
   CustomerID  Age  Gender  Tenure  Usage Frequency  Support Calls  \
0            1   22  Female     25        14             4
1            2   41  Female     28        28             7
2            3   47    Male     27        10             2
3            4   35    Male      9        12             5
4            5   53  Female     58        24             9

   Payment Delay Subscription Type Contract Length  Total Spend  \
0           27          Basic      Monthly       598
1           13         Standard      Monthly       584
2           29         Premium      Annual       757
3           17         Premium     Quarterly      232
4            2         Standard      Annual       533

   Last Interaction  Churn
0                 9     1
1                20     0
2                21     0
3                18     0
4                18     0
```

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 64374 entries, 0 to 64373
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   CustomerID      64374 non-null   int64  
 1   Age              64374 non-null   int64  
 2   Gender            64374 non-null   object 
 3   Tenure            64374 non-null   int64  
 4   Usage Frequency  64374 non-null   int64  
 5   Support Calls    64374 non-null   int64  
 6   Payment Delay    64374 non-null   int64  
 7   Subscription Type 64374 non-null   object 
 8   Contract Length  64374 non-null   object 
 9   Total Spend      64374 non-null   int64  
 10  Last Interaction 64374 non-null   int64  
 11  Churn             64374 non-null   int64  
dtypes: int64(9), object(3)
memory usage: 5.9+ MB
None
```

```
Summary Statistics:
   CustomerID      Age      Tenure  Usage Frequency  \
count  64374.000000  64374.000000  64374.000000  64374.000000
```

```

mean    32187.500000   41.970982   31.994827   15.080234
std     18583.317451   13.924911   17.098234   8.816470
min     1.000000    18.000000   1.000000   1.000000
25%    16094.250000   30.000000   18.000000   7.000000
50%    32187.500000   42.000000   33.000000   15.000000
75%    48280.750000   54.000000   47.000000   23.000000
max    64374.000000   65.000000   60.000000   30.000000
count   Support Calls  Payment Delay  Total Spend  Last Interaction \
       64374.000000  64374.000000  64374.000000  64374.000000

```

```

X = data[['Tenure']]    # independent variable
y = data['Total Spend'] # dependent variable

```

```

# Split data into 80% training and 20% testing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

```

```

# Create and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Get coefficients
print("Intercept (b0):", model.intercept_)
print("Slope (b1):", model.coef_[0])

```

```

Intercept (b0): 534.883464580102
Slope (b1): 0.175090532919694

```

```

# Predict on test data
y_pred = model.predict(X_test)

# Compare actual vs predicted
comparison = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})
print(comparison.head())

```


	Actual	Predicted
15476	635	538.385275
34666	631	539.610909
50474	314	545.388897
7984	527	543.112720
20227	236	544.688534

```

# Scatter plot with regression line
plt.figure(figsize=(8,6))
sns.scatterplot(x=X_test.squeeze(), y=y_test, color='blue', label='Actual')
sns.lineplot(x=X_test.squeeze(), y=y_pred, color='red', label='Predicted Line')
plt.title("Simple Linear Regression: Tenure vs Monthly Charges")
plt.xlabel("Tenure (Months)")
plt.ylabel("Monthly Charges")
plt.legend()
plt.show()


```

Simple Linear Regression: Tenure vs Monthly Charges



```
# Plot residuals
residuals = y_test - y_pred
plt.figure(figsize=(8,6))
sns.scatterplot(x=y_pred, y=residuals)
plt.axhline(0, color='red', linestyle='--')
plt.title("Residual Plot")
plt.xlabel("Predicted Values")
plt.ylabel("Residuals (Errors)")
plt.show()
```

Residual Plot



```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# Evaluate model performance
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print("Mean Squared Error:", mse)
print("R-squared Score:", r2)
```

```
Mean Squared Error: 68233.63038286162
R-squared Score: -0.00020630155329692812
```